Parallelization for Dummies

Graham Ganssle, Ph.D.
gra.m-gan.sl

High level computational langs

-  - lowest level here; memory leaks, heartburn, etc.
-  - huge number of ancillary functions through libs
-  - fast like C, easy (scripting) like Python

“Automatic” parallelization

- C - OpenMP / CUDA

Mature, robust platforms.
Extreme efficiency, but you’ve got to be, like, an expert, man.

- Python - joblib / DisPy

Plus a zillion other libs.
Stop working so hard. What you’re looking for exists already (it’s Python, duh).

- Julia - pmap() / @parallel

Super fun, but CPU parallel.
GPU solutions do exist; for example see the Julia CUDA package, CUDArt.jl.

AWS cloud rentals

- Cheap compute - EC2
- Repeatable builds - Vagrant
- Easy peasy storage - S3 / EBS

Requisite skills:

- Basic unix stuff
- SSH

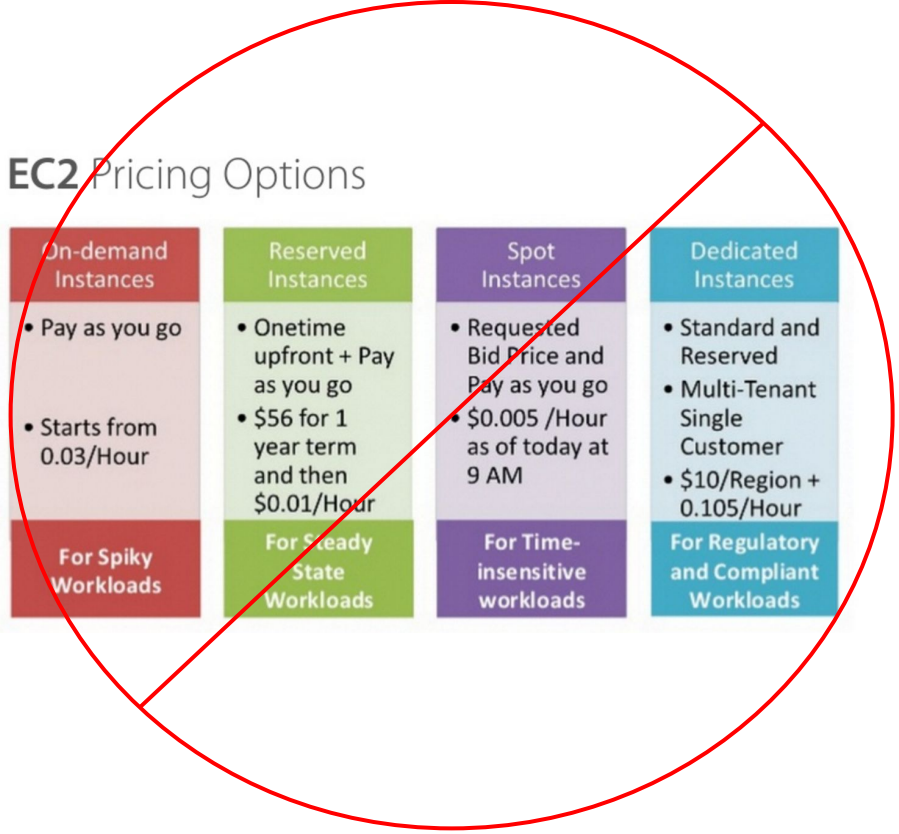

S3 ~= 3 ¢/GB

CPU vs GPU parallelization

- GPU is awsm / highly efficient
  - Have to rebuild your algorithmic flow
- CPU is less efficient, but easy
  - Operates on basically the same algo

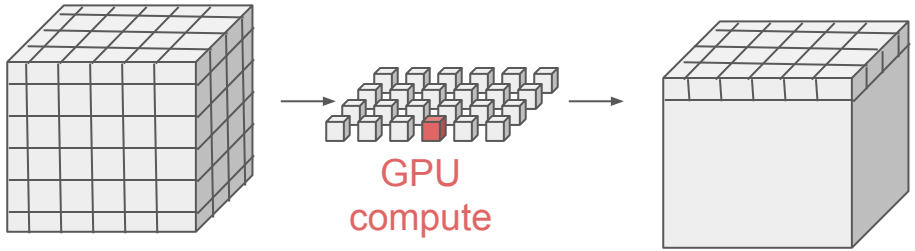

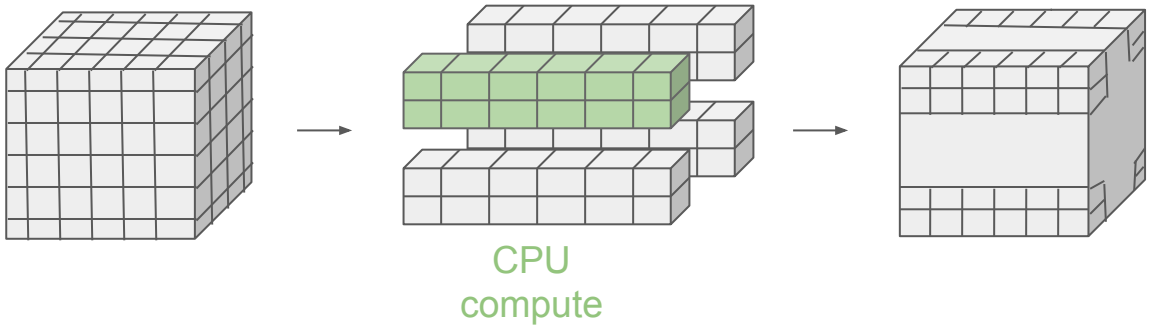

Application areas

- Inverse problems
- Optimization solutions
- Combinatorial algos

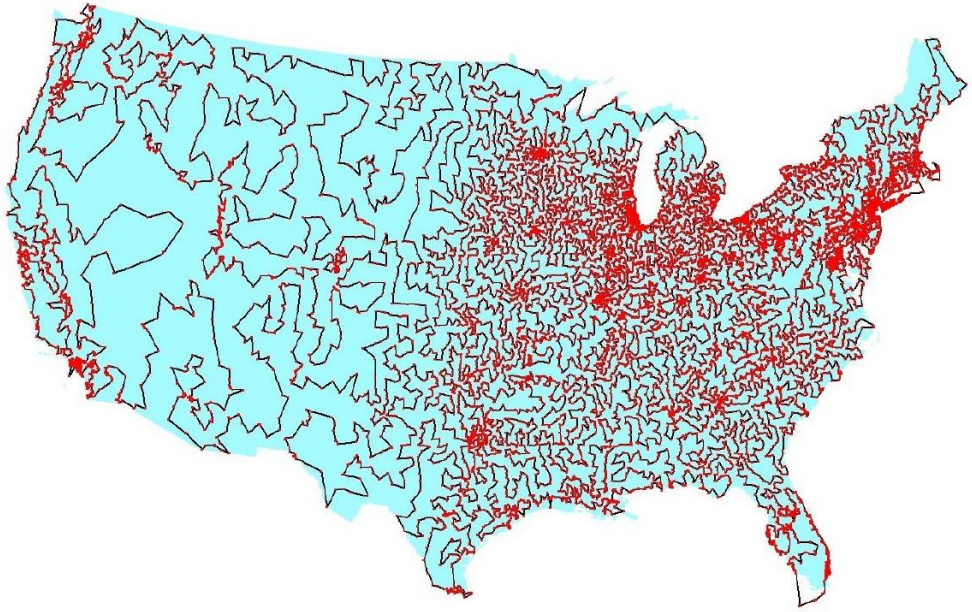

$$\frac{\partial J(\xi, \xi')}{\partial \xi} = \sum_{\gamma,\mu} S_{\mu}^{\dagger} S_{\mu} [u_{\gamma}(\xi) - u'_{\gamma}(\xi')]$$

Parallel model pays out **28%** faster than serial!

Profit!

1. Spend extra $ up front to CPU parallelize
2. Run things more quickly
3. Make money.

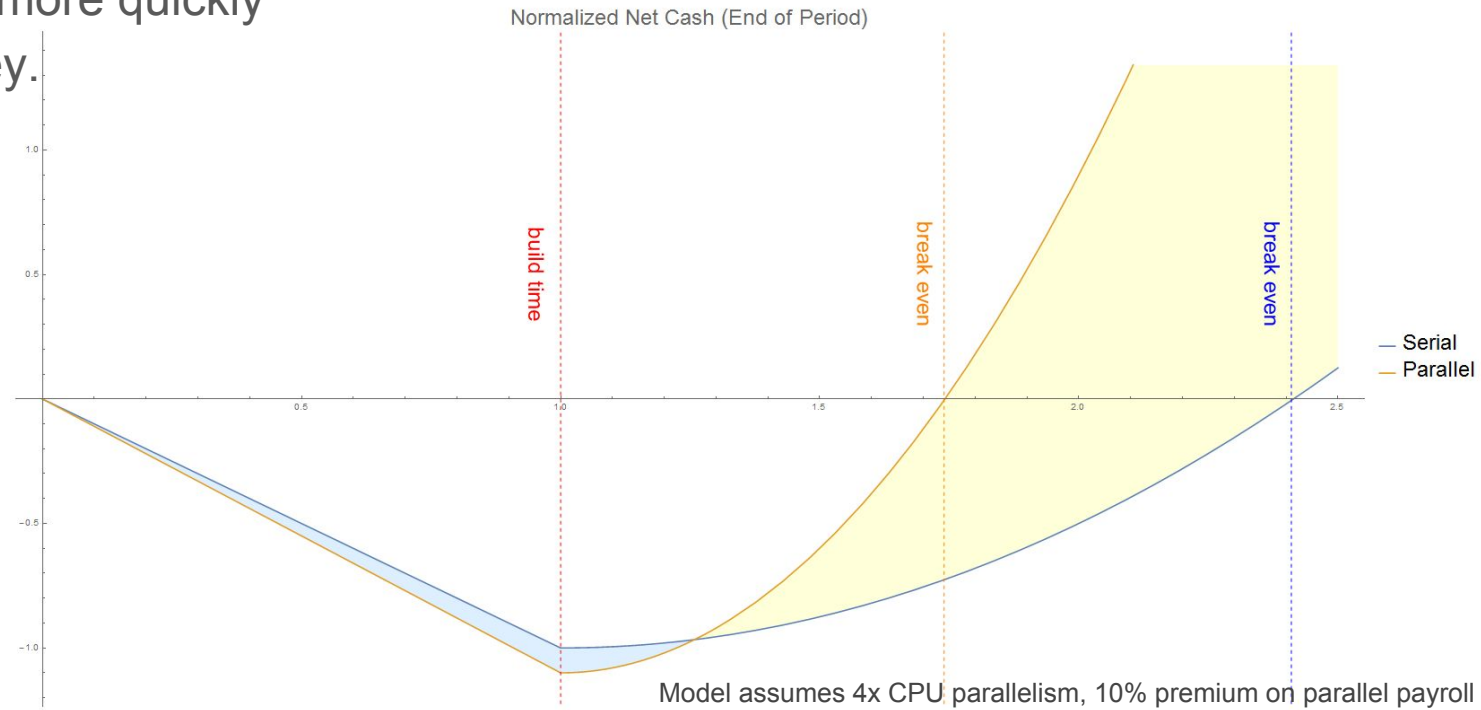

Model assumes 4x CPU parallelism, 10% premium on parallel payroll

Acknowledgements

Jaqcui Nelson

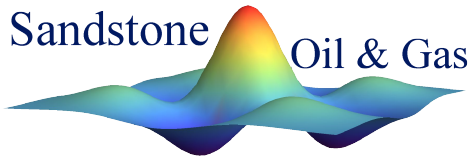

R&D DESIGN

Quickie Trial Run

- ssh -i passwd/path os@ec2-xxx.yyy.zz.qq
- s3fs {bucket} folder/ -passwd_file as/df
- Julia
    - SharedArray()
    - @sync @parallel for (i = 1; i < x; i++)
- bash time {script}